

Machine Learning Workflow

Plan

1. **Model Selection Tips** 💡
2. **Pipelines** 🔥
 - A. Preprocessing Pipes
 - Pipelines $\rightarrow \rightarrow \rightarrow$
 - Column Transformers \vdash
 - Custom Transformers \rightarrow
 - Feature Unions $\|$
 - B. Full Pipes (Preprocessing + Models)
3. **Surprise** 😲

1. Model Selection

Let's take a step back: which models have we seen so far?

1 Regression models are **parametric**

- \hat{y}
- An arbitrarily large number n of datapoints can be modeled with few β parameters

Note: Neural Networks are also parametric models (See Deep Learning)

✓ Fast to train, even on large datasets with Stochastic Gradient Descent

! Requires prior assumptions
 f

about the structure of the data; may not find complex patterns, unless given complex features

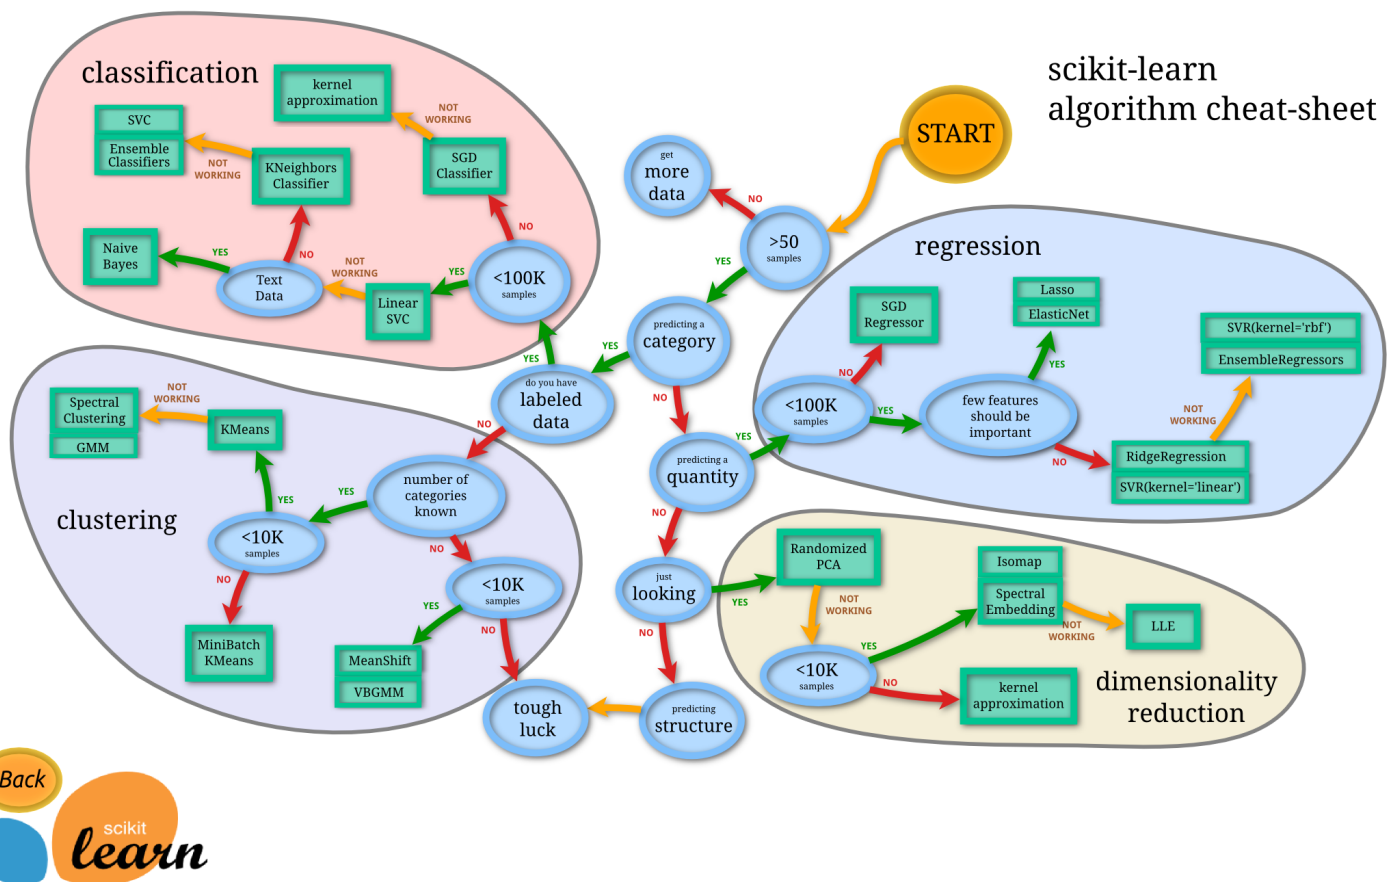
2 KNN, kernel-SVM are **non-parametric**

- No prior assumptions about the data structure are needed
- Possibly many parameters to learn (not known beforehand)
 - e.g. KNN `.fit` stores the *whole dataset*
 - e.g. rbf-SVM `.fit` must compute a Kernel between *each pair* of datapoints

Note: Trees are also non-parametric models (See Ensemble Methods)

✓ Can find complex features for you!

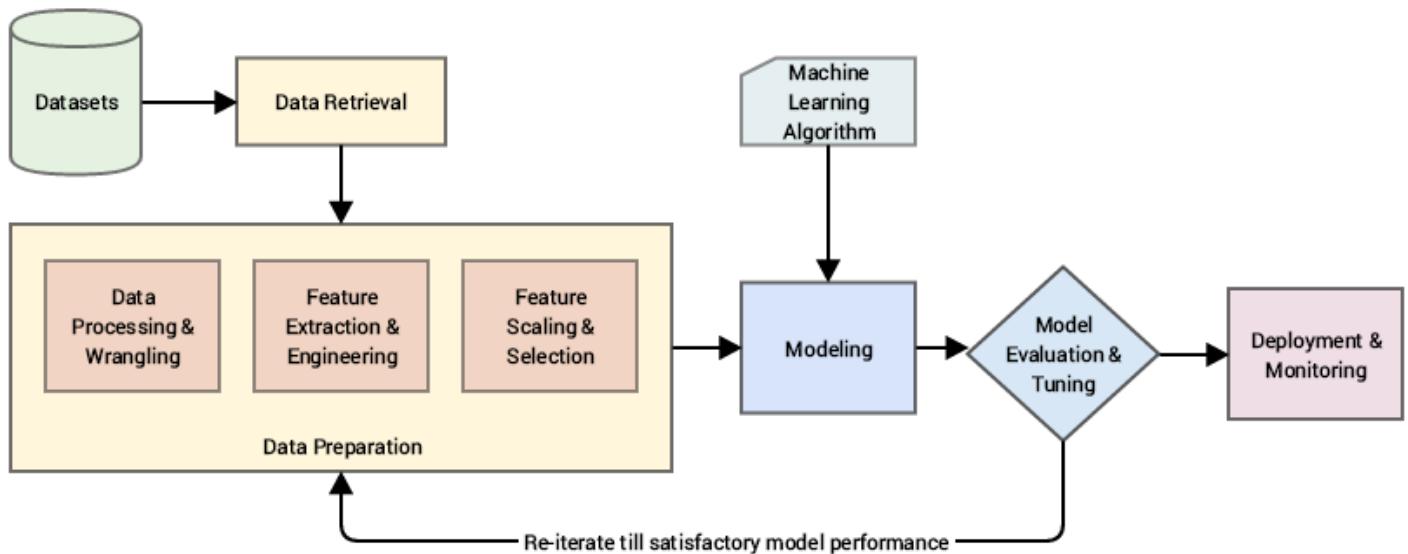
! Harder to train on large datasets and prone to overfitting



2. Pipelines




 [sklearn - Pipeline and composite estimators](https://scikit-learn.org/stable/modules/compose.html) (https://scikit-learn.org/stable/modules/compose.html)


 [sklearn.pipeline](https://scikit-learn.org/stable/modules/classes.html#module-sklearn.pipeline) (https://scikit-learn.org/stable/modules/classes.html#module-sklearn.pipeline)



A **Pipeline** is a chain of operations in a Machine Learning project (preprocessing, training, predicting, etc.)

Pipelines are powerful because they:

-  make your workflow much easier to read and understand
-  enforce the implementation and order of steps in your project
-  make your work reproducible and deployable



2.1 Preprocessing Pipelines

🎯 We are going to predict the **charges** of a health insurance contract based on various features using the following dataset.

💾 Download the dataset [here \(https://wagon-public-datasets.s3.amazonaws.com/data_workflow.csv\)](https://wagon-public-datasets.s3.amazonaws.com/data_workflow.csv)

```
In [ ]: data.head(5)
```

Out[]:

	age	bmi	children	smoker	region	charges
0	19.0	27.900	0	True	southwest	16884.92400
1	18.0	33.770	1	False	southeast	1725.55230
2	NaN	33.000	3	False	southeast	4449.46200
3	33.0	22.705	0	False	northwest	21984.47061
4	32.0	28.880	0	False	northwest	3866.85520

```
In [ ]: data.shape
```

Out[]: (1338, 6)

```
In [ ]: # Defining the features and the target

X = data.drop(columns='charges')
y = data['charges']

# Train-Test split

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size =
0.20)
X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
Out[ ]: ((1070, 5), (268, 5), (1070,), (268,))
```



Today's challenges:

1. *Impute* missing values
2. Preprocessing:
 - *Scale* numerical features
 - *Encode* categorical features
3. *Fine-tune* your ML model **and** the preprocessing steps...

... 🔥 in one cell ! 🔥

a) Pipeline → → →

A Pipeline essentially **chains** multiple steps **in sequence** (e.g. *imputing* then *scaling*)



[sklearn.pipeline.Pipeline](https://scikit-learn.org/0.16/modules/generated/sklearn.pipeline.Pipeline.html) (<https://scikit-learn.org/0.16/modules/generated/sklearn.pipeline.Pipeline.html>)

```
from sklearn.pipeline import Pipeline
```

```
In [ ]: # Preprocess "age"
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler

# Build the pipeline with the different steps
pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy="median")),
    ('standard_scaler', StandardScaler())
])

pipeline.fit(X_train[['age']])
pipeline.transform(X_train[['age']])
```


```
Out[ ]: array([[ 1.03287039],
               [-1.45497346],
               [ 1.1750329 ],
               ...,
               [ 0.25097661],
               [-0.17551091],
               [-1.2417297 ]])
```

```
In [ ]: # Show the different steps of the pipeline
pipeline
```

```
Out[ ]: Pipeline
├── SimpleImputer
└── StandardScaler
```



b) Column Transformer ♡

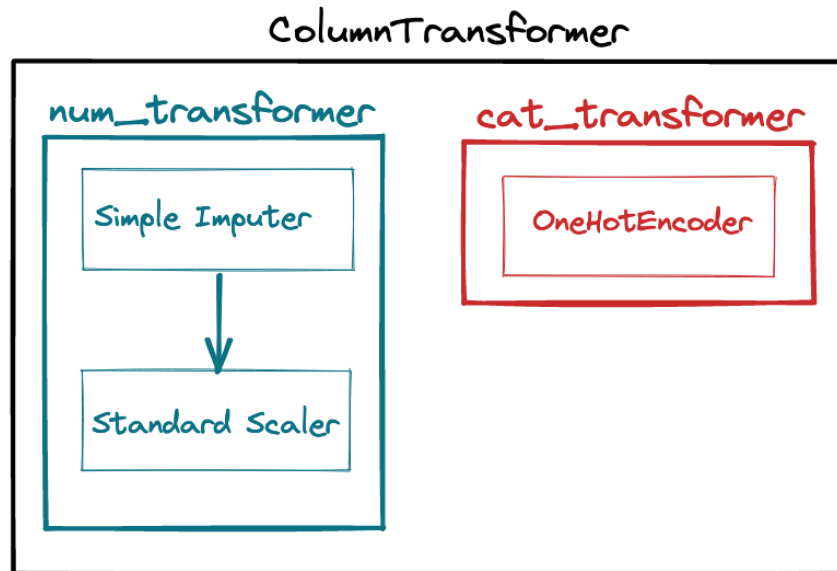
Column Transformers allow you to apply specific changes to specific columns **in parallel**

 [sklearn.compose.ColumnTransformer](https://scikit-learn.org/stable/modules/generated/sklearn.compose.ColumnTransformer.html) (<https://scikit-learn.org/stable/modules/generated/sklearn.compose.ColumnTransformer.html>)

```
from sklearn.compose import ColumnTransformer
```

Let's perform the following operations **in parallel**:

-  *Impute* then *scale* numerical values
-  *Encode* categorical values



👉 Notice how a `Pipeline` object can be passed into a `ColumnTransformer` !

```

In [ ]: from sklearn.compose import ColumnTransformer

        from sklearn.pipeline import Pipeline
        from sklearn.impute import SimpleImputer
        from sklearn.preprocessing import StandardScaler

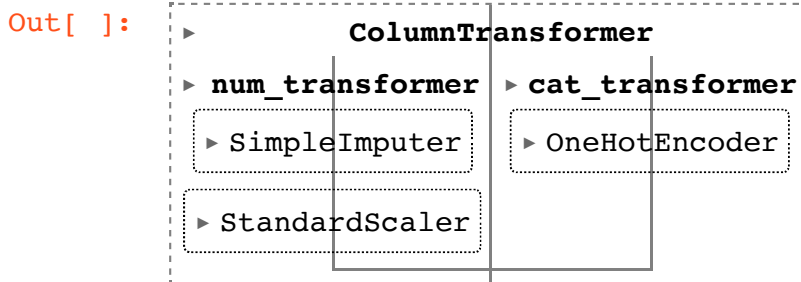
        from sklearn.preprocessing import OneHotEncoder

        # Impute then scale numerical values:
        num_transformer = Pipeline([
            ('imputer', SimpleImputer(strategy="mean")),
            ('standard_scaler', StandardScaler())
        ])

        # Encode categorical values
        cat_transformer = OneHotEncoder(handle_unknown='ignore')

        # Parallelize "num_transformer" and "cat_transformer"
        preprocessor = ColumnTransformer([
            ('num_transformer', num_transformer, ['age', 'bmi']),
            ('cat_transformer', cat_transformer, ['smoker', 'region'])
        ])
  
```

```
In [ ]: # Visualizing Pipelines in HTML
from sklearn import set_config; set_config(display='diagram')
preprocessor
```



```
In [ ]: X_train_transformed = preprocessor.fit_transform(X_train)

print("Original training set")
display(X_train.head(3))

print("Preprocessed training set")
display(pd.DataFrame(X_train_transformed).head(3))
```

Original training set

	age	bmi	children	smoker	region
162	54.0	39.60	1	False	southwest
410	19.0	17.48	0	False	northwest
639	56.0	33.66	4	False	southeast

Preprocessed training set

	0	1	2	3	4	5	6	7
0	1.032979	1.456688	1.0	0.0	0.0	0.0	0.0	1.0
1	-1.454870	-2.170790	1.0	0.0	0.0	1.0	0.0	0.0
2	1.175141	0.482582	1.0	0.0	0.0	0.0	1.0	0.0






Where are the columns' names?

😬 Don't worry and stay tuned to `scikit-learn` updates !

- scikit-learn.org/stable/whats_new.html (https://scikit-learn.org/stable/whats_new.html)

🚀 `get_feature_names_out()` 🚀

- New in `scikit-learn` 1.0.2 (September 2021)
 -  This new method helps retrieve the names of the features which went through some transformations like `StandardScaler` or `OheHotEncoder`
 -  Not all the transformers in Scikit-Learn have this new method
- New in `scikit-learn` 1.1.3: (October 2022)
 -  ALL the transformers have this method!

```
In [ ]: # Get your features' names
preprocessor.get_feature_names_out()
```

```
Out[ ]: array(['num_transformer__age', 'num_transformer__bmi',
              'cat_transformer__smoker_False', 'cat_transformer__smoker_True',
              'cat_transformer__region_northeast',
              'cat_transformer__region_northwest',
              'cat_transformer__region_southeast',
              'cat_transformer__region_southwest'], dtype=object)
```

```
In [ ]: pd.DataFrame(
        X_train_transformed,
        columns=preprocessor.get_feature_names_out()
    ).head()
```

```
Out[ ]:
```

	num_transformer__age	num_transformer__bmi	cat_transformer__smoker_False	cat_transformer__smoker_True
0	1.032979	1.456688	1.0	
1	-1.454870	-2.170790	1.0	
2	1.175141	0.482582	1.0	
3	-0.815138	0.157880	0.0	
4	-0.601894	-0.148783	0.0	

🤔 What happened to the `children` column? What if we want to keep it untouched?

👉 **remainder=passthrough**

```
In [ ]: preprocessor = ColumnTransformer([
    ('num_transformer', num_transformer, ['age', 'bmi']),
    ('cat_transformer', cat_transformer, ['region', 'smoker'])],
    remainder='passthrough')

preprocessor
```


```
Out[ ]: ColumnTransformer
├── num_transformer
│   ├── SimpleImputer
│   └── StandardScaler
├── cat_transformer
│   └── OneHotEncoder
└── remainder
    └── passthrough
```

```
In [ ]: pd.DataFrame(preprocessor.fit_transform(X_train),
    columns=preprocessor.get_feature_names_out()).head(3)
```


```
Out[ ]:
```

	num_transformer_age	num_transformer_bmi	cat_transformer_region_northeast	cat_transf
0	1.032979	1.456688	0.0	
1	-1.454870	-2.170790	0.0	
2	1.175141	0.482582	0.0	

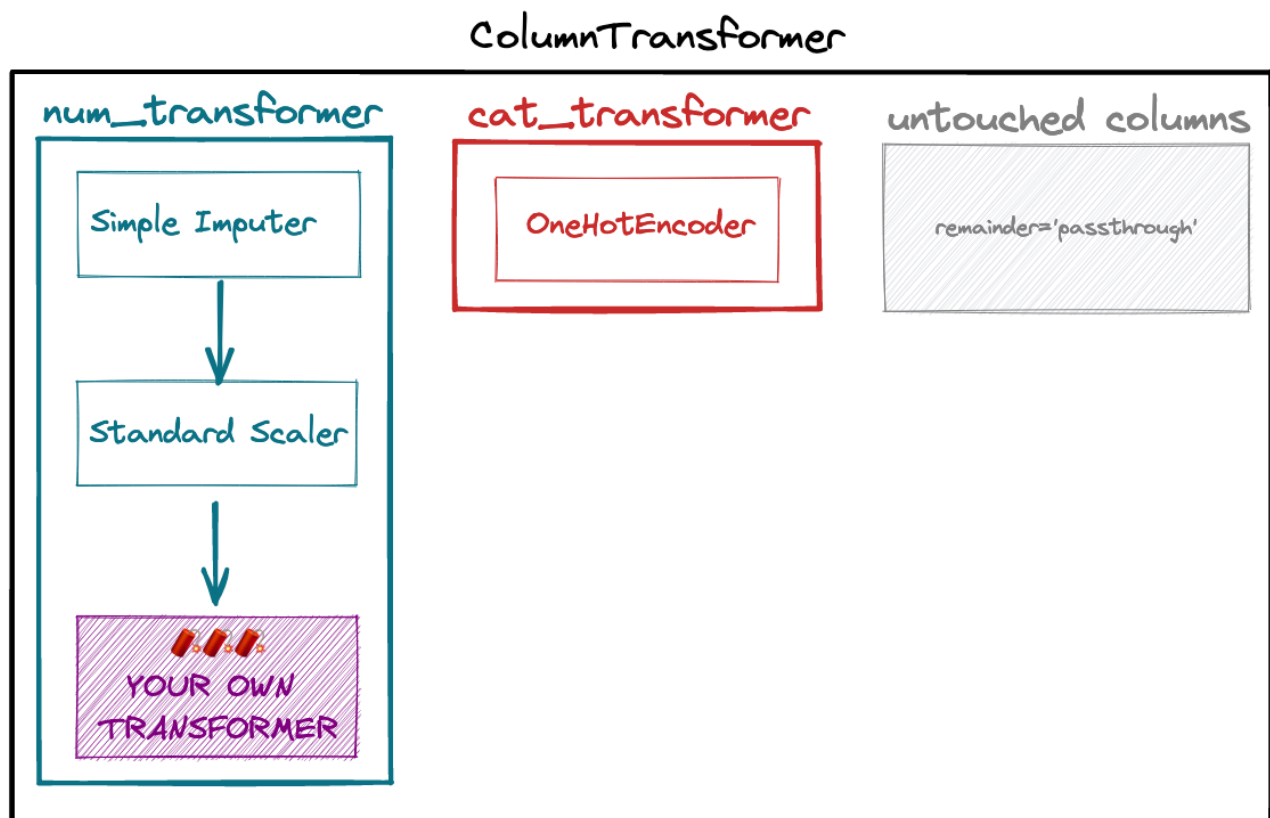
c) Custom: Function Transformer →


 [sklearn.preprocessing.FunctionTransformer](https://scikit-learn.org/stable/modules/generated/sklearn.preprocessing.FunctionTransformer.html) (<https://scikit-learn.org/stable/modules/generated/sklearn.preprocessing.FunctionTransformer.html>)

```
from sklearn.preprocessing import FunctionTransformer
```

 Function Transformers enable you to encapsulate a *Python* function within a *scikit* Transformer (→) Object

 They can be used with either Pipelines (→ → →) or ColumnTransformers (⌵)



 If you want to use your own transformer in a Pipeline or a ColumnTransformer (*not one already available in Sklearn*), you must encapsulate your function within a **FunctionTransformer**.

```
In [ ]: from sklearn.preprocessing import FunctionTransformer
```

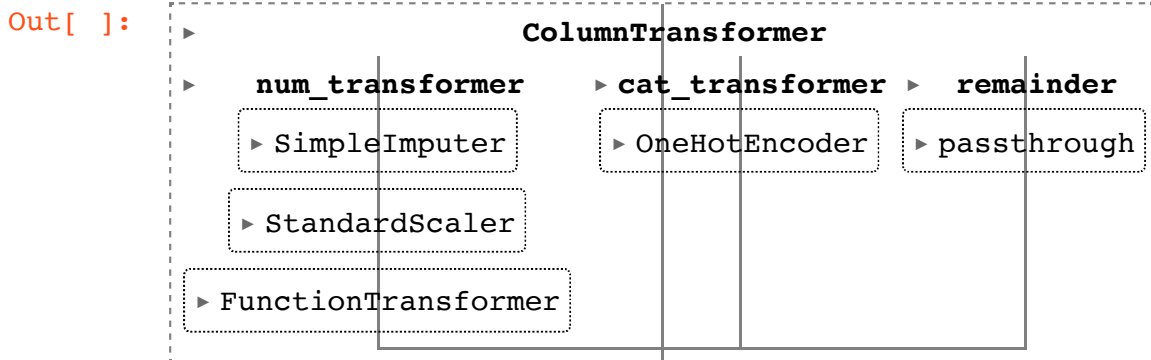
```
In [ ]: # Create a transformer that compresses data to 2 digits (for instance!)
# rounder = FunctionTransformer(np.round)

# We can use a lambda function for more customizable functions
rounder = FunctionTransformer(lambda array: np.round(array, decimals=2))
```

```
In [ ]: # Add it at the end of our numerical transformer
num_transformer = Pipeline([
    ('imputer', SimpleImputer()),
    ('scaler', StandardScaler()),
    ('rounder', rounder)])

# Encode categorical values
cat_transformer = OneHotEncoder(drop='if_binary',
                                handle_unknown='ignore')

preprocessor = ColumnTransformer([
    ('num_transformer', num_transformer, ['bmi', 'age']),
    ('cat_transformer', cat_transformer, ['region', 'smoker'])],
    remainder='passthrough')
preprocessor
```





```
In [ ]: pd.DataFrame(preprocessor.fit_transform(X_train)).head(3)
```

```
Out[ ]:
```

	0	1	2	3	4	5	6	7
0	1.46	1.03	0.0	0.0	0.0	1.0	0.0	1.0
1	-2.17	-1.45	0.0	1.0	0.0	0.0	0.0	0.0
2	0.48	1.18	0.0	0.0	1.0	0.0	0.0	4.0

! **FunctionTransformer** only works for stateless transformations !


 **stateless transformations** are transformations which cannot *store* information during `.fit(X_train)` that would be used for the `.transform(X_test)`.

 Since a stateless transformation doesn't learn anything, fitting it is impossible, it does nothing other than transform!

 `FunctionTransformer` is compatible with stateless transformations.

Examples of transformations which don't "learn" anything:

- $X \rightarrow \log(X)$
- $(X_1, X_2) \rightarrow X_1 + 5X_2$

 **stateful transformations** are transformations which *store* information during `.fit(X_train)`. This information is re-used for `.transform(X_test)`.

Examples of transformations which "learn" something:

$X_{train} \rightarrow StandardScaler(X_{train})$

learns

μ_{train}

and

σ_{train}

$X_{train} \rightarrow MinMaxScaler(X_{train})$

learns

$X_{train}^{(min)}$

and

$X_{train}^{(max)}$

 `FunctionTransformer` is not compatible with stateful transformations

 We will have to code our own `Class` to use `FunctionTransformer` with stateful transformations!

 **Transformers under the hood**

```
In [ ]: from sklearn.base import TransformerMixin, BaseEstimator



class MyCustomTranformer(TransformerMixin, BaseEstimator):
    # BaseEstimator generates the get_params() and set_params() methods
    # that all Pipelines require
    # TransformerMixin creates the fit_transform() method from fit() and
    # transform()

    def __init__(self):
        pass

    def fit(self, X, y=None):
        # Here you store what needs to be stored/learned during .fit(X
        # _train) as instance attributes
        # Return "self" to allow chaining .fit().transform()
        pass


    def transform(self, X, y=None):
        # Return the result as a DataFrame for an integration into the
        # ColumnTransformer
        pass
```

```
my_transformer = MyCustomTranformer()
my_transformer.fit(X_train)
my_transformer.transform(X_train)
my_transformer.transform(X_test)
```

 More in today's challenges 

d) FeatureUnion ||

`FeatureUnion` applies a list of transformer objects **in parallel** to the input data, then **concatenates** the results. This is useful to combine several feature extraction mechanisms into a single transformer

 [sklearn.pipeline.FeatureUnion \(https://scikit-learn.org/stable/modules/generated/sklearn.pipeline.FeatureUnion.html\)](https://scikit-learn.org/stable/modules/generated/sklearn.pipeline.FeatureUnion.html)

 Useful to **create entirely new features!**

Example: let's build and add a new feature called `bmi_age_ratio`

```
In [ ]: X_train.head(3)
```

Out[]:

	age	bmi	children	smoker	region
162	54.0	39.60	1	False	southwest
410	19.0	17.48	0	False	northwest
639	56.0	33.66	4	False	southeast

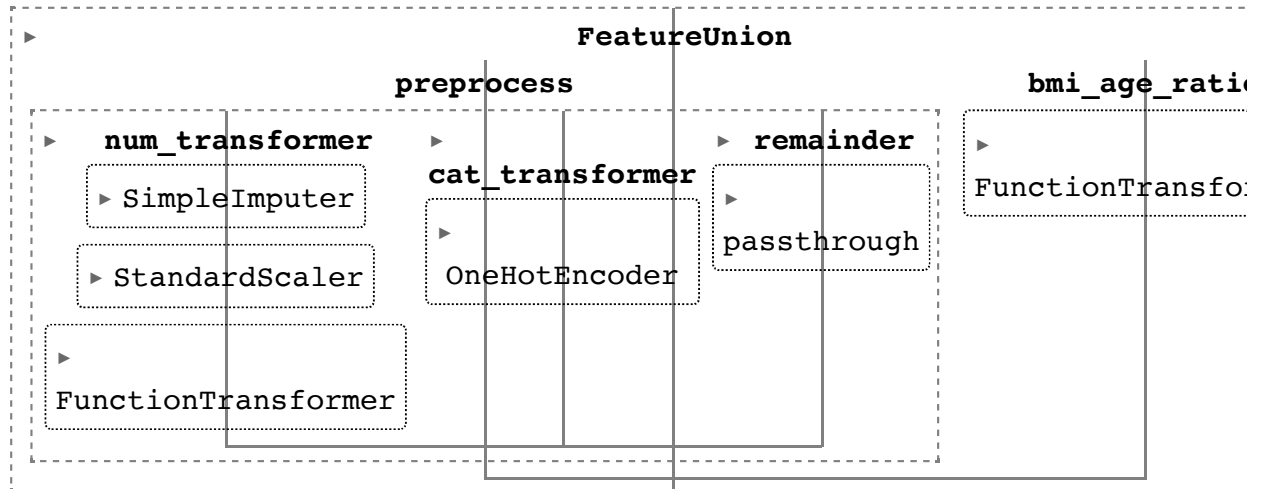
```
In [ ]: from sklearn.pipeline import FeatureUnion

# Create a custom transformer that multiplies/divides two columns
# Notice that we are creating this new feature completely randomly just as an example
bmi_age_ratio_constructor = FunctionTransformer(lambda df: pd.DataFrame(df["bmi"] / df["age"]))

union = FeatureUnion([
    ('preprocess', preprocessor), # columns 0-7
    ('bmi_age_ratio', bmi_age_ratio_constructor) # new column 8
])

union
```

Out[]:



```
In [ ]: pd.DataFrame(union.fit_transform(X_train)).head(1)
```

```
Out[ ]:
```

	0	1	2	3	4	5	6	7	8
0	1.46	1.03	0.0	0.0	0.0	1.0	0.0	1.0	0.733333

Building your preprocessor with `make_***` shortcuts ⚡

```
from sklearn.pipeline import Pipeline
from sklearn.pipeline import FeatureUnion
from sklearn.compose import ColumnTransformer
```

There are equivalent transformers using the syntax `make_***` 📌


```
In [ ]: from sklearn.pipeline import make_pipeline
        from sklearn.pipeline import make_union
        from sklearn.compose import make_column_transformer
```

```
Pipeline([
    ('my_name_for_the_imputer', SimpleImputer()),
    ('my_name_for_the_scaler', StandardScaler())
])
```

⇔

```
In [ ]: make_pipeline(SimpleImputer(), StandardScaler())
```

```
Out[ ]:
```

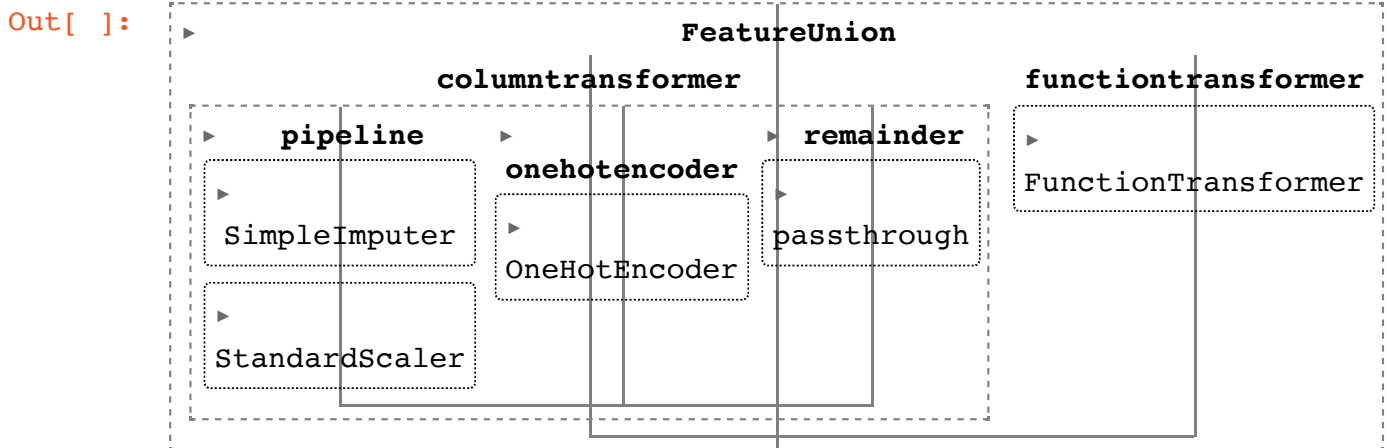



```
In [ ]: num_transformer = make_pipeline(SimpleImputer(), StandardScaler())
cat_transformer = OneHotEncoder()

preproc_basic = make_column_transformer(
    (num_transformer, ['age', 'bmi']),
    (cat_transformer, ['smoker', 'region']),
    remainder='passthrough'
)

preproc_full = make_union(preproc_basic, bmi_age_ratio_constructor)

preproc_full
```



 **make_column_selector** selects features automatically based on dtype

```
from sklearn.compose import make_column_selector
```

```
num_col = make_column_selector(dtype_include=['float64'])
cat_col = make_column_selector(dtype_include=['object', 'bool'])
```

```
In [ ]: X_train.dtypes
```

```
Out[ ]: age          float64
bmi           float64
children      int64
smoker        bool
region        object
dtype: object
```

🎨 Complete preprocessing pipeline 🎨

```
In [ ]: from sklearn.compose import make_column_selector

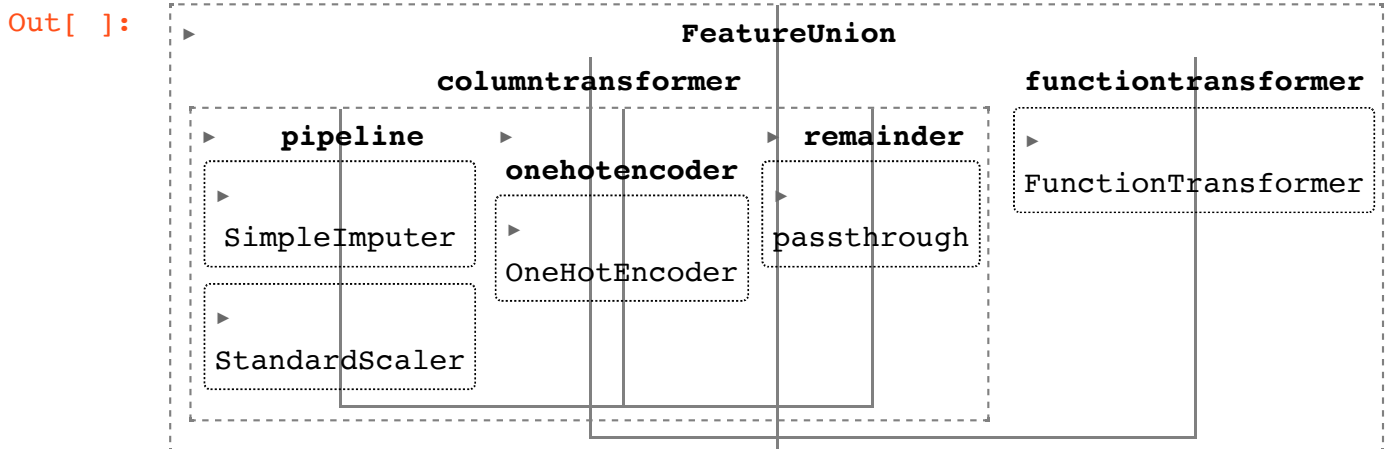
num_transformer = make_pipeline(SimpleImputer(), StandardScaler())
num_col = make_column_selector(dtype_include=['float64'])

cat_transformer = OneHotEncoder()
cat_col = make_column_selector(dtype_include=['object', 'bool'])

preproc_basic = make_column_transformer(
    (num_transformer, num_col),
    (cat_transformer, cat_col),
    remainder='passthrough'
)

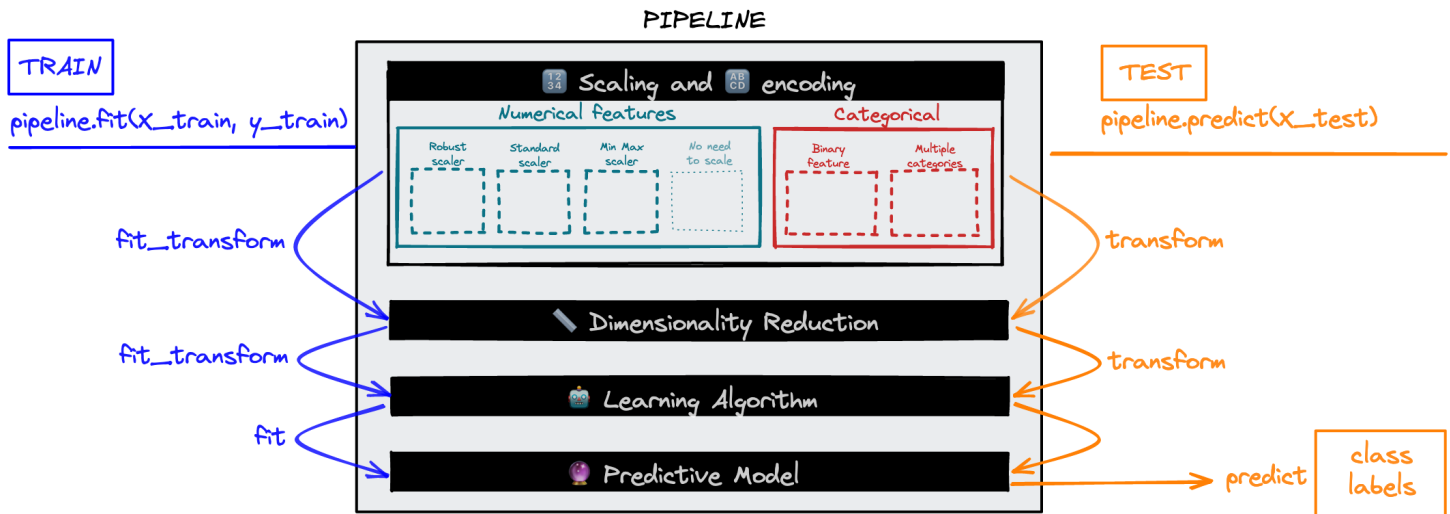
preproc_full = make_union(preproc_basic, bmi_age_ratio_constructor)

preproc_full
```



2.2 Including models in Pipelines

- Model objects can be plugged into Pipelines
- Pipelines inherit the methods of the **last** object in the sequence
 - Transformers: `fit` and `transform`
 - Models: `fit`, `score`, `predict`, etc.



- When executing the `pipeline.fit` method, the transformer's `.fit` and `.transform` methods will be called sequentially, and the model will be trained.
 - At this stage, all transformers' variables are saved into the memory of the pipeline
- When executing the `pipeline.predict` method, only the transformer's `.transform` method will be called, using the variables learned during the original `fit`

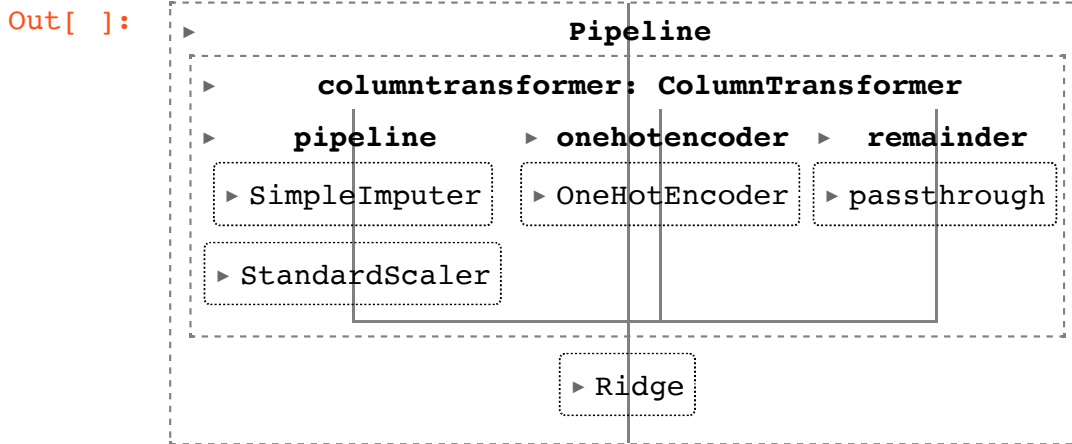
a) Full pipeline

```
In [ ]: from sklearn.linear_model import Ridge

# Preprocessor
num_transformer = make_pipeline(SimpleImputer(), StandardScaler())
cat_transformer = OneHotEncoder()

preproc = make_column_transformer(
    (num_transformer, make_column_selector(dtype_include=['float64', 'float32'])),
    (cat_transformer, make_column_selector(dtype_include=['object', 'boolean'])),
    remainder='passthrough'
)

# Add estimator
pipeline = make_pipeline(preproc, Ridge())
pipeline
```



```
In [ ]: # Train Pipeline
pipeline.fit(X_train,y_train)

# Make predictions
pipeline.predict(X_test.iloc[0:1])

# Score model
pipeline.score(X_test,y_test)
```

Out[]: 0.7473478157212925

b) Cross-validate a Pipeline

```
In [ ]: from sklearn.model_selection import cross_val_score

# Cross-validate Pipeline
cross_val_score(pipeline, X_train, y_train, cv=5, scoring='r2').mean()

Out[ ]: 0.7434317676218065
```

c) Grid Search a Pipeline

- *Grid Searching* allows you to check which combination of preprocessing/modeling **hyperparameters** works best.
- It is possible to *Grid Search* the hyperparameters of **any component of the Pipeline**
 - Typical Sklearn syntax: `step_name__transformer_name__hyperparameter_name`
 - To check which hyperparameters of the pipeline can be optimized: `pipeline.get_params()`

```
In [ ]: # Which parameters of the pipeline are GridSearch-able?
pipeline.get_params()
```

```
In [ ]: from sklearn.model_selection import GridSearchCV

grid_search = GridSearchCV(
    pipeline,
    param_grid={
        # Access any component of the Pipeline
        # and any available hyperparameter you want to optimize
        'columntransformer__pipeline__simpleimputer__strategy': ['mean', 'median'],
        'ridge__alpha': [0.1, 0.5, 1, 5, 10]
    },
    cv=5,
    scoring="r2")

grid_search.fit(X_train, y_train)

grid_search.best_params_
```

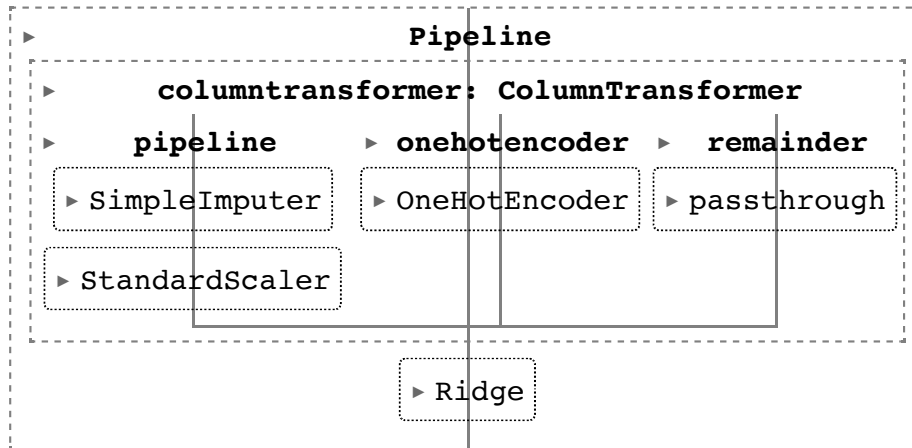
```
Out[ ]: {'columntransformer__pipeline__simpleimputer__strategy': 'mean',
        'ridge__alpha': 1}
```




Let's save the pipelined model with the best hyperparameters.

```
In [ ]: pipeline_tuned = grid_search.best_estimator_  
pipeline_tuned
```

```
Out[ ]:
```



 We can use this "best" model for predictions without re-training it!

```
In [ ]: pipeline_tuned.predict(X_test[0:1])
```

```
Out[ ]: array([10216.56989159])
```

d) Caching to avoid repeated computations

🤔 Are your preprocessing steps too long to run?

🔑 You can use caching techniques!

```
from tempfile import mkdtemp
from shutil import rmtree

# Create a temp folder
cachedir = mkdtemp()

# Instantiate the Pipeline with the cache parameter
pipeline = Pipeline(steps, memory=cachedir)

# Clear the cache directory after the cross-validation
rmtree(cachedir)
```

With the parameter `memory=cachedir`, `preproc` parameters can be cached into memory.

- Avoid recalculating all of the parameters during `CrossValidation` or `GridSearchCV` on `estimator` hyperparams only
- Helpful only when the transformer's `.fit` time is long and the dataset is very large

e) Debug your pipe

```
In [ ]: # Access the components of a Pipeline with `named_steps`
        pipeline_tuned.named_steps.keys()
```

```
Out[ ]: dict_keys(['columntransformer', 'ridge'])
```

```
In [ ]: # Check intermediate steps
        print("Before preprocessing, X_train.shape = ")
        print(X_train.shape)
        print("After preprocessing, X_train_preprocessed.shape = ")
        pipeline_tuned.named_steps["columntransformer"].fit_transform(X_train).shape
```

```
Before preprocessing, X_train.shape =
(1070, 5)
After preprocessing, X_train_preprocessed.shape =
```

```
Out[ ]: (1070, 9)
```

f) Exporting models/Pipelines



You can export your final model/pipeline as a `pickle` file



The file can then be loaded back into a notebook or deployed on a server (see `ML_Ops` module).

```
In [ ]: import pickle

# Export Pipeline as pickle file
with open("pipeline.pkl", "wb") as file:
    pickle.dump(pipeline_tuned, file)

# Load Pipeline from pickle file
my_pipeline = pickle.load(open("pipeline.pkl", "rb"))

my_pipeline.score(X_test, y_test)
```

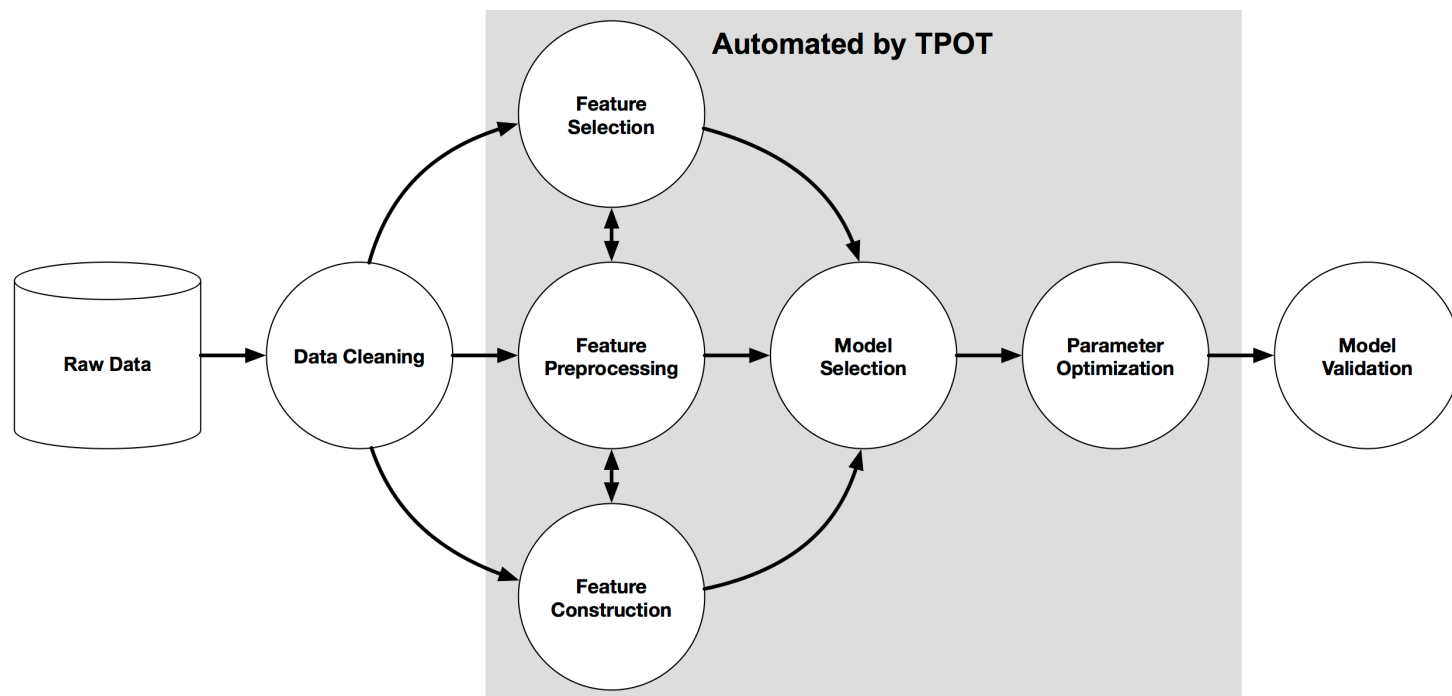
```
Out[ ]: 0.7473478157212925
```

3. Surprise

AutoML


TPOT

The Tree-based Pipeline Optimization Tool (TPOT) is an automated Machine Learning tool that optimizes Machine Learning Pipelines



 More details available in the [TPOT documentation \(http://epistasislab.github.io/tpot/\)](http://epistasislab.github.io/tpot/)

Installation



```
pip install TPOT
```

```
In [ ]: import os
        from tpot import TPOTRegressor

        X_train_preproc = preproc_basic.fit_transform(X_train)
        X_test_preproc = preproc_basic.transform(X_test)
```

```
In [ ]: # Instantiate TPOTClassifier
tpot = TPOTRegressor(generations=4, population_size=20, verbosity=2, scoring='r2', n_jobs=-1, cv=2)

# Process autoML with TPOT
tpot.fit(X_train_preproc, y_train)

# Print score
print(tpot.score(X_test_preproc, y_test))
```

Generation 1 - Current best internal CV score: 0.8517440046999218

Generation 2 - Current best internal CV score: 0.853008927910814

Generation 3 - Current best internal CV score: 0.853008927910814

Generation 4 - Current best internal CV score: 0.8558530771102855

Best pipeline: RidgeCV(GradientBoostingRegressor(input_matrix, alpha=0.9, learning_rate=0.01, loss=ls, max_depth=3, max_features=0.6000000000000001, min_samples_leaf=14, min_samples_split=13, n_estimators=100, subsample=0.55))
0.872811877434264

/Users/davywai/.pyenv/versions/3.8.12/envs/lewagon-data/lib/python3.8/site-packages/sklearn/metrics/_scorer.py:765: FutureWarning: sklearn.metrics.SCORERS is deprecated and will be removed in v1.3. Please use sklearn.metrics.get_scorer_names to get a list of available scorers and sklearn.metrics.get_metric to get scorer.
warnings.warn(

```
In [ ]: # Export TPOT Pipeline to a Python file
tpot.export(os.path.join(os.getcwd(), 'tpot_iris_pipeline.py'))

! cat 'tpot_iris_pipeline.py'

import numpy as np
import pandas as pd
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.linear_model import RidgeCV
from sklearn.model_selection import train_test_split
from sklearn.pipeline import make_pipeline, make_union
from tpot.builtins import StackingEstimator

# NOTE: Make sure that the outcome column is labeled 'target' in the
data file
tpot_data = pd.read_csv('PATH/TO/DATA/FILE', sep='COLUMN_SEPARATOR',
dtype=np.float64)
features = tpot_data.drop('target', axis=1)
training_features, testing_features, training_target, testing_target
= \
    train_test_split(features, tpot_data['target'], random_s
tate=None)

# Average CV score on the training set was: 0.8558530771102855
exported_pipeline = make_pipeline(
    StackingEstimator(estimator=GradientBoostingRegressor(alpha=0.9,
learning_rate=0.01, loss="ls", max_depth=3, max_features=0.600000000
0000001, min_samples_leaf=14, min_samples_split=13, n_estimators=10
0, subsample=0.55)),
    RidgeCV()
)

exported_pipeline.fit(training_features, training_target)
results = exported_pipeline.predict(testing_features)
```


Summary

- Pipeline
→
list of sequential steps

- `ColumnTransformer`
→
list of parallel steps
 - `remainder="passthrough"` : used to save untransformed columns
- `FunctionTransformer`
→
encapsulates a function as a Scikit-Learn transformer that you can plug into a Pipeline or a `ColumnTransformer`
- `FeatureUnion`
→
applies transformations in parallel and concatenates the results, quite useful for feature creation

```
In [ ]: preproc_full
```

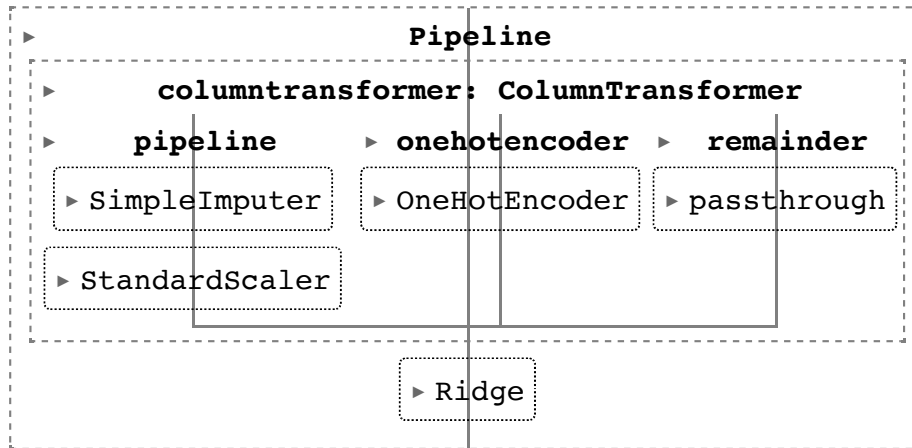
```
Out[ ]:
```



- You can chain a `preprocessor pipeline` with a Scikit Learn model
- A full pipeline can go through `cross_validate`, `GridSearchCV`, `RandomizedSearchCV`

In []: pipeline_tuned

Out[]:



Your turn! 🚀