



ML Ops Module

 [lecture starter for teacher](#)

 **Welcome to WagonCab** 

 You've just been hired by WagonCab as a **Machine Learning Engineer**

But what is an ML Engineer, exactly?

 It's a new role that is emerging as data jobs **specialize**

Data Scientist

Data Engineer



Core Competencies
Adv. Math/Statistics
ML/AI
Adv. Analytics



Overlapping Skills
Analysis
Programming
Big Data



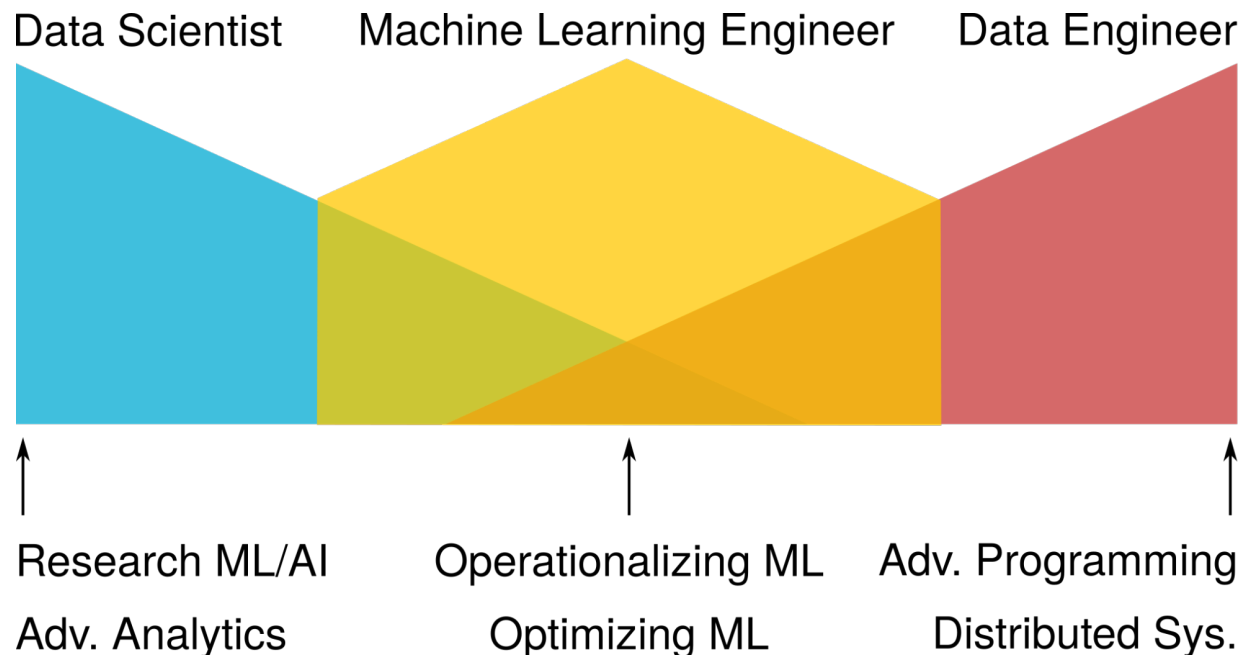
Core Competencies
Adv. Programming
Distributed Sys.
Data Pipelines

BDI

BIG DATA INSTITUTE

For more information go to <http://bigdatainstitute.io>

Attribution-NoDerivatives 4.0 International (CC BY-ND 4.0) Licensed



BDI

BIG DATA INSTITUTE

For more information go to <http://bigdatainstitute.io>

Attribution-NoDerivatives 4.0 International (CC BY-ND 4.0) Licensed

💡 It's also emerging as **tools** become more mature

🧐 The full modern ML stack is way too big to cover in this bootcamp ([source](#))



🚗 **Back to WagonCab** 🚗

Your company is launching a new ML product called **TaxiFare**

🎯 Learn to predict the price of conventional taxi rides in New York

🎯 Show in-app → 📱 how much users would save using **WagonCab** 🚗 instead of Taxis 🚕!



VS



WagonCab has the huge, public [NYC Trip Record Dataset](#) at its disposal, which is around 170GB in size, and looks as follows

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
0	4.5	2009-06-15 17:26:21 UTC	-73.8443	40.7213	-73.8416	40.7123	1
1	16.9	2010-01-05 16:52:16 UTC	-74.016	40.7113	-73.9793	40.782	1
2	5.7	2011-08-18 00:35:00 UTC	-73.9827	40.7613	-73.9912	40.7506	2
3	7.7	2012-04-21 04:30:42 UTC	-73.9871	40.7331	-73.9916	40.7581	1

A team of **Data Scientists** has already created an ML model to predict the price of a ride 🤖

However:

- Work has been done in a isolated context (Notebook)

- The dataset used to train was a smaller, more manageable subset of the NYC dataset (100K rows)

🎯 Your goal as an **ML engineer** will be to:

- train the model at scale 💪
- train the model in the cloud ☁️
- deploy the model in production 🚢
- manage model lifecycle (performance monitoring, re-training on new data, etc.) 🔁
- provide a user interface to access it 🎨

Unit 1) Train at Scale

🎯 Today's goals:

- Understand data scientists' notebooks
- **Package** your Python code (👤 lecture)
- Master your **IDE** (👤 lecture)
- Learn **incremental** processing techniques to handle GBs worth of data

1) Packaging & Virtual Env 101

1.1) What is a Package?

👉 Re-usable code from one project to another (`from ... import ...`)

A package allows you to:

👉 Share it with others

- Install from [PyPI](#): `pip install <package_name>`
- Install from [GitHub](#): `pip install git+https://...`

👉 Deploy in production (on Linux servers)

👉 Track code (git) and collaborate on it!

🎯 **Lecture's goal:** create a package called `toto` that you will be able to install on any machine
`pip install toto`

Anatomy of a Minimal Python Package

```

.                                # project directory
├─ setup.py                     # lists package name and dependencies
└─ toto                         # package directory
    ├─ __init__.py              # defines toto as a package
    └─ lib.py                   # your code

```

- A **module** is a single python file inside a package
- A **package** is directory of python modules that contains an additional `__init__.py`

👉 `__init__.py` allows you to write `toto.lib`, for instance

CLI

`python -m toto.lib` # *executes it as module*

`python toto/lib.py` # *executes it as file*

python file

`from toto.lib import a_function`

👉 `__init__.py`'s content (often empty) is executed at each import line

🖥️ **LIVECODE: minimal package**

`mkdir project-toto`

`cd project-toto`

`mkdir toto`

`touch toto/lib.py`

`touch toto/__init__.py`

`touch setup.py`

`code .`

`toto/lib.py`

```

def who_am_i():
    print("Hello my name is Jean")

```

```

if __name__ == '__main__':
    who_am_i()

```

1.2) Virtual Environments

🖥️ Let's create a dedicated virtual env `totoenv` for this project

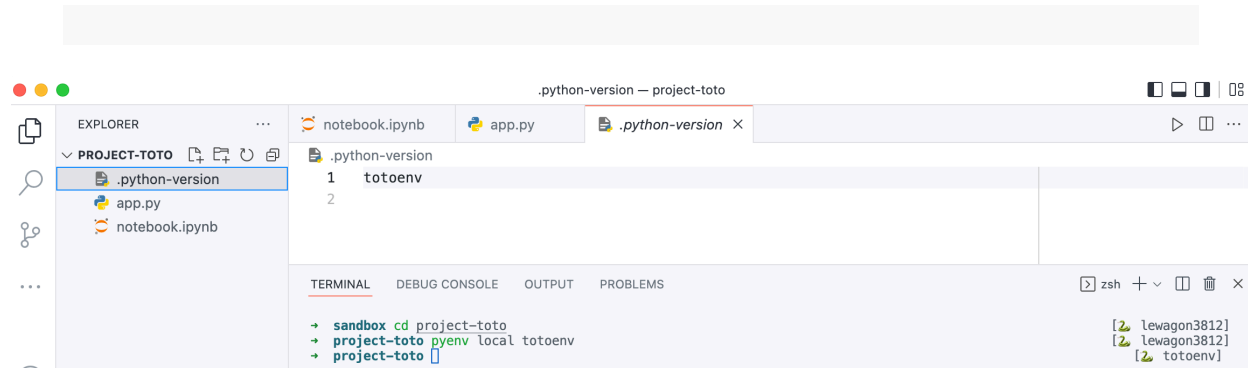
💡 One `venv` per project is a good practice!

Create new totoenv inside python 3.8.12

`pyenv virtualenv 3.8.12 totoenv`

In project-toto, create `.python-version` that activates totoenv when present

`pyenv local totoenv`



Install minimal packages for this demo lecture

`pip list`

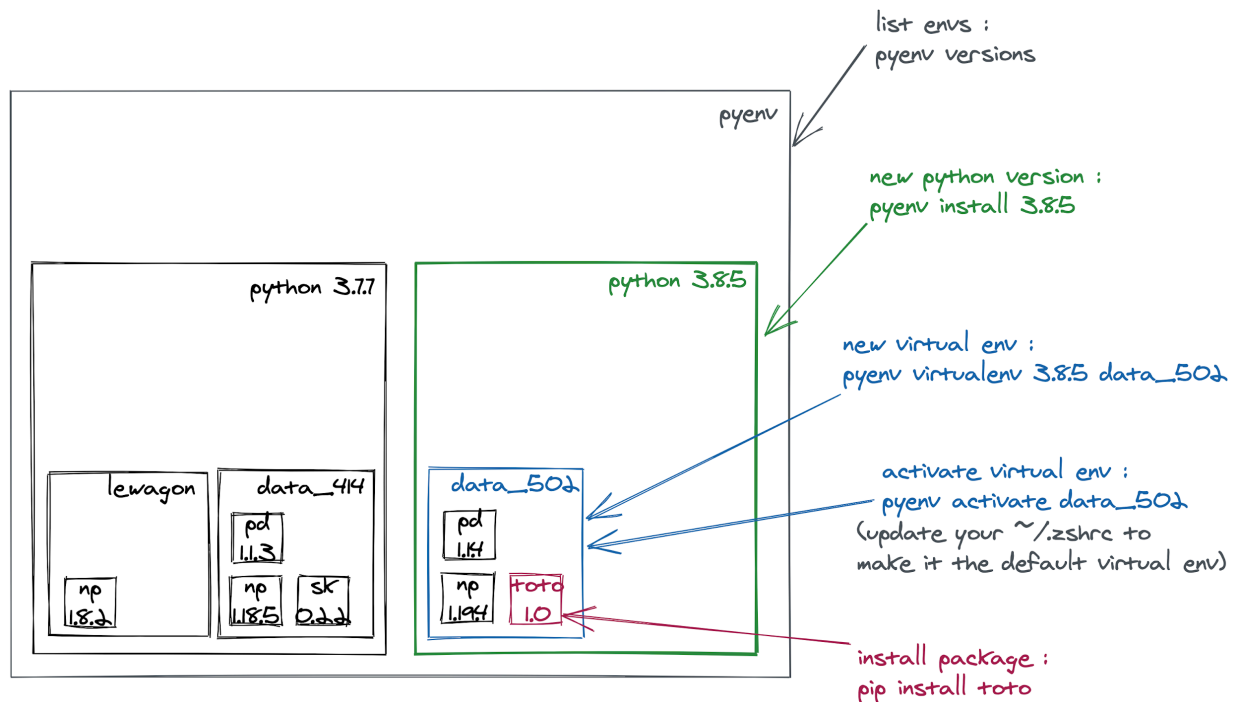
`pip install --upgrade pip`

`pip install pandas`

`pip install ipython` *# needed for ipython*

`pip install ipykernel` *# needed for notebooks*

Reminder on pyenv vs. venv



2) 🖥 Installing & Using a Package

🎯 Goals

- `pip install toto` in virtual env `totoenv`

2.1) Install the Package

Fill in `setup.py`

`setup.py`

from `setuptools` import `setup`

```
setup(name='toto',  
      description="package description",  
      packages=["toto"]) # You can have several packages, try it
```

Install

Sit next to `setup.py` and run:

`pip install .`

Verify that the package is installed

`pip freeze`

pip freeze | grep toto

pip freeze G toto

this is an Oh-My-Zsh shortcut

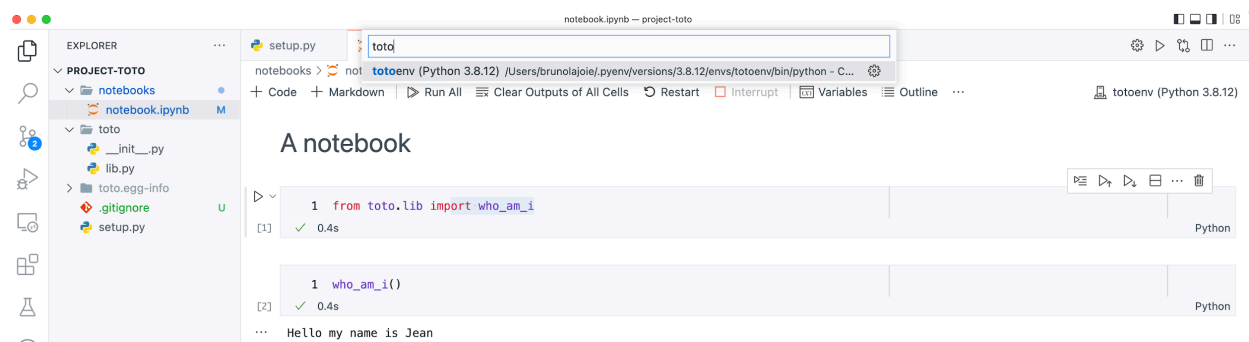
2.2) Run the Package from Anywhere (when **totoenv** is Activated)

For instance, from a notebook

mkdir notebooks

touch notebooks/notebook.ipynb

Open the notebook with VS Code, select **ipykernel=totoenv**, and you should be able to call **from toto.lib import who_am_i**
who_am_i()



? Why Does it Work?

Well, remember that Python's **import** always looks at your **PYTHONPATH**

import sys

sys.path

Which pyenv always appends with **site-packages**:

👉 `~/pyenv/versions/3.8.12/envs/totoenv/lib/python3.8/site-packages`

🔍 ...and where is **toto** located?

import toto

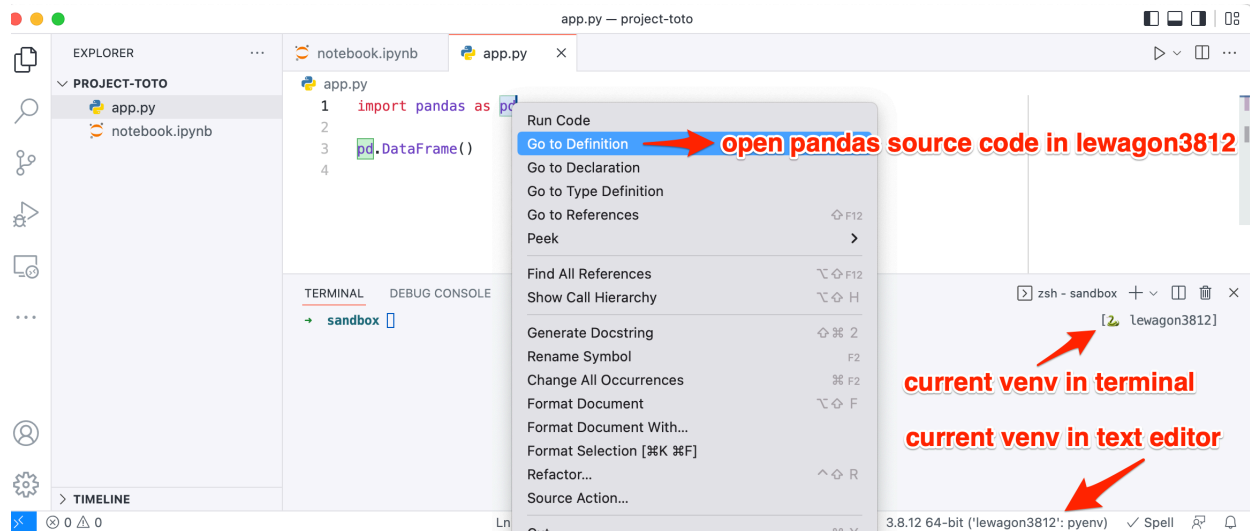
toto.__file__

👉 `~/pyenv/versions/3.8.12/envs/totoenv/lib/python3.8/site-packages/toto/__init__.py`

✅ **pip install .** creates a **COPY** of your project folder inside **site-packages**

`site-packages` contains all your pip packages.

🔍 Use it to explore third-party libraries!



🔥 Hot-Reload on the Package?

```
def who_am_i():
```

```
    print("Hello my name is Jean UPDATED")
```

👉 Is the modification visible?

❌ No, because `pip install .` only creates a **COPY** of your project in `site-packages`

✅ `pip install -e .` (editable) for hot-reloading

- First, `pip uninstall toto` (move away from `setup.py`'s root to do it properly)
- Then `pip install -e .`

```
import toto
```

```
toto.__file__ # '~/code/sandbox/project-toto/toto/__init__.py'
```

PS.: for notebooks and `ipython`, don't forget the magic command to avoid closing/reopening

```
ipython every time
```

```
%load_ext autoreload
```

```
%autoreload 2
```

🔥 **Pro Tip:** setup `autoreload` as default action

```
//settings.json
```

```
"jupyter.runStartupCommands": [  
    "%load_ext autoreload",  
    "%autoreload 2"  
],
```

2.3) Dependencies

👉 Let's say we want the `termcolor` package to be installed along with `toto`

```
# toto/lib.py
```

```
from termcolor import colored
```

```
def who_am_i():  
    print(colored("Hello my name is Jean", "blue"))
```

✅ **Solution:** create `requirements.txt` and update `setup.py`

```
# Terminal
```

```
touch requirements.txt
```

```
echo termcolor >> requirements.txt
```

...or specify the version to be used (`termcolor==1.1.0`, for example)

```
# setup.py
```

```
from setuptools import setup  
from setuptools import find_packages
```

```
# list dependencies from file
```

```
with open('requirements.txt') as f:  
    content = f.readlines()  
requirements = [x.strip() for x in content]
```

```
setup(name='toto',  
      description="package description",  
      packages=find_packages(), # NEW: find packages automatically  
      install_requires=requirements) # NEW
```

🚧 Then, `pip install -e .` to update `totoenv` with `requirements.txt`

```

.
├─ notebooks
|   └─ notebook.ipynb      # your code
├─ toto                    # package directory
|   └─ __init__.py         # defines toto as a package
|   └─ lib.py              # your code
├─ requirements.txt        # lists the dependencies
├─ setup.py               # lists package name and dependencies
├─ .python-version        # stores name of virtual env
└─ .gitignore             # files not to track with git

```

2.4) Adding a **Makefile** to Create Simple CLI Commands ⚡

📝 syntax: `make <some_action>`

```

.
├─ notebooks
|   └─ notebook.ipynb      # your code
├─ toto                    # package directory
|   └─ __init__.py         # defines toto as a package
|   └─ lib.py              # your code
├─ Makefile               # command line directive manager
├─ requirements.txt        # lists the dependencies
├─ setup.py               # lists package name and dependencies
├─ .python-version        # stores name of virtual env
└─ .gitignore             # files not to track with git

```

Makefile Sample

Makefile

directive_name:

<tab>some command with all its arguments

<tab>@this command will not print out before being executed

<tab>-the command after this one will run no matter what

<tab>-@the markers can be combined

A Simple Makefile

Makefile

install:

```
@pip install -e .
```

clean:

```
@rm -f */version.txt
```

```
@rm -f .coverage
```

```
@rm -f */ipynb_checkpoints
```

```
@rm -Rf build
```

```
@rm -Rf */__pycache__
```

```
@rm -Rf */*.pyc
```

all: install clean

Let's run a directive; while sitting next to the **Makefile**:

```
tree
```

```
make clean
```

```
tree
```

Let's run several directives; while sitting next to the **Makefile**:

```
make install clean # = make all
```

! The **Makefile** is highly sensitive:

- Its name is case-sensitive; call it **makefile** and nothing works
- The commands inside of the directives must be indented exclusively using **tabulations**; use one or more spaces instead and nothing works

3) Testing your Package 🧪

Why should we bother with tests?

- ensures robustness of the project in case of changes
- allows teams to work on the same project without breaking each other's code
- with code, you can "describe" what your code should do better than with words 🚫 (TDD)

👉 **Test Driven Development** (TDD) consists of writing the tests **before** the actual code

👉 Most software teams hire full-time testers!

Test Example

toto/divide.py

```
def divide_without_raising(x:float, y:float) -> float:
    """
    divides x by y, but instead of raising errors when y equals 0, returns:
    - inf if x positive
    - -inf if x negative
    - nan if x equals 0
    """
    pass # YOUR CODE HERE
```

💡 Small parentheses regarding `inf` and `nan`

```
inf = float('inf')
assert type(inf) == float
```

We can do arithmetic on infinity!

```
assert inf == inf
assert inf + inf == inf
assert inf * inf == inf
assert inf * -2 == -inf
```

BUT these operations on inf are undefined ❌

```
inf - inf
inf / inf
```

nan

```
nan = float('nan')
```

```
assert type(nan) == float
assert nan != nan # nan is the ONLY FLOAT that does NOT equal itself
```

import math

```
assert math.isnan(nan) # check for "nanism"
assert math.isnan(nan + 2) # ANY operation on nan is nan
```

Let's do **TDD**: write tests before coding the function

import math

from toto.divide import divide_without_raising

```
def test_has_correct_arithmetic():
    assert divide_without_raising(2.0, 2.0) == 1.0, 'wrong basic arithmetic'
```

```
def test_handles_divide_by_zero_correctly():
```

```
assert divide_without_raising(2., 0.) == float('inf')
assert divide_without_raising(-2., 0.) == -1 * float('inf')
assert math.isnan(divide_without_raising(0., 0.))
```

Launch your tests using the **pytest** framework 🔥

```
echo pytest >> requirements.txt
pip install -e .
pytest tests -v # verbose
```

You'll often see it written in the **Makefile** so you can **make test** your package

```
# Makefile
test:
    @pytest -v tests
```

Write all your tests the same place

```
.
├─ notebooks
│   └─ notebook.ipynb    # your code
├─ tests
│   └─ test_lib.py       # a test file for lib.py
├─ toto                  # package directory
│   └─ __init__.py       # defines toto as a package
│   └─ lib.py            # your code
├─ Makefile              # command line directive manager
├─ requirements.txt      # lists the dependencies
├─ setup.py              # lists package name and dependencies
├─ .python-version       # stores name of virtual env
└─ .gitignore            # files not to track with git
```

SOLUTION

```
def divide_without_raising(x:float, y:float) -> float:
    """
    divides x by y, but instead of raising errors when y equals 0, returns:
    - inf if x positive
    - -inf if x negative
```

```

- nan if x equals 0
"""
if y != 0.:
    return x/y
else:
    if x > 0.:
        return float('inf')
    if x < 0.:
        return -1 * float('inf')
    if x == 0.:
        return float('nan')

```

🚩 Lastly: add a **README.md** to help reproduce your work!

```

README.md
# How to install
pip install toto

# How to reproduce results
from toto.lib import who_am_i
who_am_i()

# How to run tests
make tests

```

4) Data Engineering Tips 💡

4.1) Become a Debugging Master!

💻 Live Demo:

- add call to `who_am_i()` inside `divide_without_raising()`
- change color from "blue" to ""

Learn how to read your **stack trace**

Traceback (most recent call last):

```

File "toto/divide.py", line 22, in <module>
    divide_without_raising(2.,0.)
File "toto/divide.py", line 10, in divide_without_raising
    who_am_i()
File "/Users/brunolajoie/code/sandbox/project-toto/toto/lib.py", line 5, in who_am_i
    print(colored("Hello my name is Jean", ""))
File "/Users/brunolajoie/.pyenv/versions/totoenv/lib/python3.8/site-packages/termcolor.py", line 105,
in colored

```



```
text = fmt_str % (COLORS[color], text)
KeyError: "
```

🔥 **Pro Tip:** use `Option-Click` to navigate to the line

🔥 **Pro Tip 2:** use `ipdb.set_trace()` (`<=>` `breakpoint()`) instead of `print()`

`pip install ipdb` (do not add it to `requirements.txt`, it's a *dev-only* package)

ipdb navigation

- `s` (step into)
- `n` (next = step over)
- `c` (continue to next error or `ipdb.set_trace()`)
- `u` (up stack trace)
- `d` (down stack trace)
- `return` (continue until current function's `return`)
- `l` (provide more context)
- `ll` (provide a lot more context)
- `q` (quit; `exit` also works)

🔥 **Pro Tip 3:** automatically set a trace where your code stopped

```
if __name__ == '__main__':
    try:
        divide_without_raising(2., 0.)
    except:
        import ipdb, traceback, sys

        exctype, value, tb = sys.exc_info()
        traceback.print_exc()
        ipdb.post_mortem(tb)
```

👉 Then just use `u` or `d` to get up/down the stack trace until you find your codebase!

4.2) Master your IDE

VS Code Shortcuts (macOS)

- Command palette `⌘-⇧-P`
- Toggle Terminal: `Ctrl-Backtick`
- Split screens with `Ctrl-⌘-⇧`
- Toggle file bar `⌘-B`
- Move panel position (palette)
- Go to file `⌘-P`
- Navigate to symbols globally `⌘-⇧-R`

- Navigate file-to-file `⌘-Click`
- Search `⌘-⇧-F`
- Replace `⌘-⇧-H`
- Rename symbols across all your files (Right-Click)
- Create your own shortcuts (palette)
- Create your own snippets (palette)
- Learn your `settings.json`

Notebooks

- magic commands `# %%` to have Jupyter-like code cells in any Python file
- `convert` from `.ipynb` to `.py` file
- setup `autoreload` as default action


//settings.json

```
"jupyter.runStartupCommands": [
  "%load_ext autoreload",
  "%autoreload 2"
],
```

4.3) Master your Command Line

 `manual`

- `man git` full manual for Git CLI

 `--help`: shorter, and works for sub-commands

- `git --help`
- `git pull --help`

 `tldr` is an even shorter summary

- `brew install tldr`
- `tldr git pull`

Aliases & Commands:

- Customize your aliases in `code ~/.aliases`
- `alias hi='echo hello world'`
- `which hi` lists the location of a command or the command behind an alias

Your Turn!

You are an **ML Engineer** at `WagonCab` now, working on putting the `TaxiFare` price predictor in production!



Challenge of the Day

- Understand Data Scientists' notebooks
- Package their code into a Python package
- Train it at scale with incremental processing techniques

Appendix

(No live lecture)

A.1) Memory Optimization

For large dataset, it may be useful to compress data by "downcasting" dtypes to **the smallest possible values** according to existing ranges in your dataset

```
s_int = pd.Series([1, 2, 134])
```

```
s_int
```

```
0    1
1     2
2   134
```

```
dtype: int64
```

[pd.Series.astype](#) allows to specify a particular [numpy data type](#)

```
s_int = s_int.astype(np.int16)
```

```
s_int
```

```
0    1
1     2
2   134
```

```
dtype: int16
```

👉 **Beware of the haircut!**

```
s_int = s_int.astype(np.int8)
```

```
s_int
```

```
0    1
1     2
2  -122
```

```
dtype: int8
```

What happened?

`np.int8` uses 8 bits to represent numbers, of which one bit is used for the sign. So it can only handle integers between -128 and 127.

Let's see what happens once you go above 127.

```
edge_number = pd.Series([127], dtype=np.int8)
```

```
edge_number
```

```
0    127
```

```
dtype: int8
```

```
edge_number + 1  # This overflows
```

```
0   -128
```

```
dtype: int8
```

So we have to make sure that our range fits within the `np.int8` range before downcasting.

```
134 < 2**7  # This number does not fit within the np.int8 range
```

False

Floats

```
s_float = np.array([0.1234567890123456789, 2, 3], dtype=np.float16)
print("16bit: ", a[0])
```

```
s_float = np.array([0.1234567890123456789, 2, 3], dtype=np.float32)
print("32bit: ", b[0])
```

```
s_float = np.array([0.1234567890123456789, 2, 3], dtype=np.float64)
print("64bit: ", c[0])
```

👉 Be careful when playing with float precision, especially when **scaling**

👉 `float32` should be enough for most of your first Data Science projects

👉 `float16` is extremely uncommon

Downcast made easy with `pd.to_numeric` 🤖

- Converts to the *smallest possible* `int` data type that do not change values
- Also turns `floats64` to `floats32`

```
s_int = pd.Series([1, 2, 134])
s_float = pd.Series([0.1234567890123456789, 2, 3], dtype=np.float64)
```

```
print(pd.to_numeric(s_int, downcast='integer'), '\n')
print(pd.to_numeric(s_float, downcast='float'))
```

```
0    1
1    2
2   134
dtype: int16
```

```
0    0.123457
1    2.000000
2    3.000000
dtype: float32
```

Assess memory usage

```
df = pd.DataFrame(dict(
    my_int=[123 for _ in range(100)], # int64
    my_float=[1.0 for _ in range(100)], # float64
    my_bool=[True for _ in range(100)]) # bool
```

df

```
df.memory_usage()
```

📁 Keep that for later use 📁

```
def compress(df, **kwargs):
```

```
    """
```

```
    Reduces size of dataframe by downcasting numerical columns
```

```
    """
```

```
    input_size = df.memory_usage(index=True).sum() / 1024
```

```
    print("new dataframe size: ", round(input_size, 2), 'kB')
```

```
    in_size = df.memory_usage(index=True).sum()
```

```
    for type in ["float", "integer"]:
```

```
        l_cols = list(df.select_dtypes(include=type))
```

```
        for col in l_cols:
```

```
            df[col] = pd.to_numeric(df[col], downcast=type)
```

```
    out_size = df.memory_usage(index=True).sum()
```

```
    ratio = (1 - round(out_size / in_size, 2)) * 100
```

```
    print("optimized size by {} %".format(round(ratio, 2)))
```

```
    print("new dataframe size: ", round(out_size / 1024, 2), " kB")
```

```
    return df
```

```
compress(df)
```

```
df.memory_usage()
```

```
df.dtypes
```

A.2) Scripts

```
📄 syntax: do_something
```

A **script** is an executable file that can run from **anywhere** on your terminal

📁 Let's add our script in **scripts/toto-run**

```
#!/usr/bin/env python
```

```
from toto.lib import who_am_i
```

```
who_am_i()
```

👉 The *shebang* `#!` line indicates what interpreter to use

Tell the package to deploy the script in **setup.py**

```
setup(name='toto',  
      ...  
      scripts=['scripts/toto-run']) # NEW LINE
```

Then `pip install -e .` again

Use script anywhere (if `toto-env` is activated)

```
toto-run
```